A Parking Guidance and Information System for TinyOS

CSE 521S Final Report

Matthew Lindsay   Patrick McBryde   Michael Schultz

Washington University in Saint Louis

ABSTRACT

The Internet is able to provide up to the minute updates on traffic conditions and provide motorists with detailed mapping and directions for nearly any destination. However, once a driver arrives at their destination, they must still search for a parking space. This wastes fuel and time, as well as causing traffic buildup and contributing to the urban heat island effect. We believe that Wireless Sensor Networks (WSN) can be used to aide in these problems. Drivers would be able to access up to the minute information about parking from their home computer or smart phone. Our system can provide high frequency updates for, we believe, a cost lower than current implementations. Most commonly, expensive induction loops or other, more permanent systems must be installed, and only provide entrance and exit information for parking structure.

Modern "Intelligent Transportation Systems" (ITS) are a mix of systems that gather various data metrics from transportation areas (parking lots, streets, alleyways, etc.), aggregate this data together, and present coherent information to the end-users of the system. Orchestrating such a system presents several challenges at each level. How do you gather the data and at what granularity? Where does the data go once gathered and can it be useful to anyone? If it can be useful, how can it be presented to end-users to provide accurate and simple-to-interpret content? This paper presents and explains our decisions in developing a "Parking Guidance and Information" (PGI) system for TinyOS.

1. INTRODUCTION

When Tim Berners-Lee introduced the World Wide Web to the world in 1991, he commented that,

> "[The World Wide Web] could start a revolution in information access."

In the past 20 years there has been a huge wealth of information placed online for anyone to access anywhere. The revolution is just beginning however, there is still an incredible amount of "now" information that simply isn't available online yet. We want to use wireless sensor networks in order to push more "now" information online, in this project we target parking lot data to simplify a driver's life when traveling.

Driving to new destinations always brings some level of stress and uncertainty. Questions like: where am I going, what do the roads and intersections look like, what side of the road is the place on, and where do I park are common when traveling to new locations. Luckily, mapping technology brings us the answers or at least partial answers to some of these questions. In the past decade there has been huge growth of web-based mapping technology, allowing these questions to be answered even more fully. For example, Google Maps allows you to get directions from point A to point B, and in the past 5 years has introduced Street View [13] to their interface. Street View allows a user to view the roads they will be traveling on from an in-car perspective. This removes much of the uncertainty when driving, leading to less confusion, a better experience, and potentially fewer accidents.

However, there is still more that can be done. Once you've arrived at your destination you must find where to park. This can be a hassle for those unfamiliar with an area or during peak hours when there may not be a parking space nearby. Parking Guidance and Information Systems (PGIS) allow drivers to quickly evaluate where empty parking spaces may exist [12]. Traditional PGI Systems just count the number of cars that enter and exit a designated area, displaying the number of spots available in that area to a driver. These systems are imprecise, costly, and are difficult to integrate with other technologies.

This paper presents our experiences building a PGI System on the TelosB/TinyOS platform. By using motes, it is possible to deploy a single sensor for each parking space allowing more detailed information about an entire parking lot. Moreover, with proper sheltering, these sensors can be mounted on the surface of a parking lot instead of cutting into the lot and installing an inductive loop to detect vehicle presence. We can also take advantage of the wireless multi-hop routing abilities of TinyOS-based motes to avoid costly wiring and enable a heterogeneous mix of sensors. Combining these aspects gives a system is able to be more precise, lower cost, and easily integrated with future technologies.

The remainder of this paper is as follows. Section 2 defines the goals of this project more fully. Section 3 presents and explains the high-level architecture of our

PGI System and details the hardware and software used during the course of this project. Section 4 talks about how we convinced ourselves that the system worked correctly and, if we have more resources, could scale up as needed. Sections 5 and 6 give an overview of related works and lessons we learned during this project. Finally, Section 7 concludes this paper with a discussion of potential future work for this project.

2. GOALS

Since we believe other solutions to PGI Systems do not represent what modern technology is capable of, our goals for this project are to build a PGI System that can be easily deployed at a low-cost while aggregating the data to a single location and integrating it with user-facing technology to provide a rich user experience. In short, we plan to answer the following questions:

- How do you gather parking data and at what granularity?

- Where does the data go once gathered and can it be useful to anyone?

- How can it be presented to end-users to provide accurate and simple-to-interpret content?

Our goals for the hardware involve prototyping a low-cost, wireless system that functions with a heterogeneous sensor suite. This can be accomplished by using sensors with standard connection ports, allowing the system to be customized with different sensing methods. We can also take advantage of the built-in sensors to provide drivers with advanced feedback, such as headlight detection. Lastly, our project should be available for current parking structures to use. Therefore, it must make use of the wireless capabilities of the mote platform, reducing installation costs by removing the need for costly wiring.

Once our system is installed, we can collect the data at a base station, located at the parking structure, which will interact with both the WSN and the Internet. Communication between the base station and the network will rely upon Collection Tree Protocol (CTP), a robust method of point-to-sink wireless communication. Our base station then communicates with a computer via USB and sends the data to the Internet.

3. DESIGN

Figure 1 shows a high-level view of the infrastructure of our PGI System. Our system begins with a parking lot with each space equipped by a sensor built on the TelosB/TinyOS platform. Data is collected at a base station and then sent over the Internet to our backend. When our backend receives the data it updates and logs the changes to the data store and organizes the lot information. Once the data is stored, clients can query the backend, getting up to the minute information about a specific lot and lots in the vicinity.

To present the system, we first introduce our API that ties the two major components (sensing and aggrega-
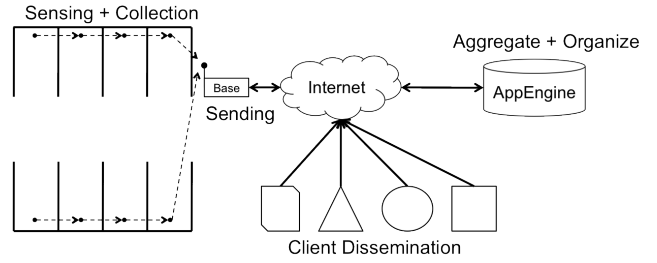
Figure 1: High-level view of our PGI System. A base station at the lot sends data over the Internet to our aggregator, which end-users can query.

```
[
  {
    "space_id": 8,
    "is_empty": false,
    "magnet": 218,
    "sonar": 114
  },
  ...
]
```

Figure 2: Client to server JSON data.

tion) together. Once the API is in place, we explain the hardware sensing design and then the aggregation and dissemination system.

3.1 A RESTful API

REpresentational State Transfer (REST) is an abstract model for building large-scale web services [11]. The principles of a RESTful architecture are an identification of resources through a URI, a "pure" HTTP interface, and self-contained hyperlinks. This section describes how we achieved these principles.

Our URI scheme is straight-forward. Each lot is given a unique identification consisting of alphanumerics and the underscore (_) character. When attempting to access a specific lot, that identifier is part of the URI (i.e., `http://<host>/lot/wustl_millbrook/` identifies the Milbrook lot on Washington University's campus).

Requests concerning this lot must go through its respective URI. Requests use one of the HTTP verbs of `GET`, `PUT`, or `POST`. A `GET` request simply returns the information for the requested URI in either the HTML or JSON [9] format. A `PUT` request attempts to store updated status information to the lot. This request contains the list of parking spaces to update, their identification and status, as well as any associated meta-information that should be put in the datastore, an example of this can be seen in Figure 2. Finally, a `POST` request handles the creation of new parking lots.

As mentioned, a `GET` request on the parking lot identifier results in either a HTML or JSON response. An

```
{
  "lot_id": "wustl_snowway",
  "geo_pt": "38.650233,-90.313529",
  "timestamp": 1291946695.0
  "spaces":
  [
    {
      "space_id": "4",
      "is_empty": false,
      "timestamp": 1291946695.0,
      "extra_info": "sonar:55;magnet:201;"
    },
    ...
  ]
}
```

Figure 3: Server to client JSON data.

example of the JSON response can be seen in Figure 3. The JSON response provides enough information for a client application to identify the name of the lot, its location, the time of the most recent activity, and the list of spaces. Each space has a unique identifier, the status of that spot (full or empty), the time it was last updated, and any additional information about that spot (likely to be sensor readings at the last report). A client can potentially combine this information with a virtual representation of the lot's topography, then keep that representation updated by periodically querying the server for new information.

With the API in place, we are able to tie an arbitrary sensor network into our service (if it follows our description) and can provide rich applications access to our information with ease and the possibility of expansion. We are now free to discuss the specifics of our hardware sensing platform and our web application that presents the data.

3.2 Hardware
When choosing the hardware for our PGI system we had specific goals in mind. The system must be low-cost, reliable, power efficient, and easily customized to meet the needs of the specific structure or parking lot. A primary focus of this system is to upgrade existing parking structures, not just add it to new construction. Therefore our system must be capable of customizing sensors for different environments.

For example, an induction loop sensor is typically the best sensor for detecting the presence of a vehicle, but in a pre-existing structure you must cut into the floor of every parking space to add them, a costly procedure. When adding an inductive loop system, the installation cost is more prohibitive than that of the PGI system. This customization is likely the most important requirement for the successful adoption of our system by the parking industry. By using wireless technology and supporting many different sensor types and configurations we believe we can successfully implement this system in both pre-existing and new parking structures.

Figure 4: Telos Revision B.

Figure 5: LV-MaxSonar®-EZ1™ by MaxBotix® (hereafter "EZ1").

3.2.1 Parking Space Monitors
A Parking Space Monitor (PSM) must be able to do more than just determine a vehicle's presence. It must also run on batteries for an extended period of time, support multiple sensor types and interfaces, and transmit the required data. For all these reasons we chose the Telos Revision B (TelosB) wireless sensor module as the base for our PSMs.

The TelosB (Figure 4) provides many features while still managing to use little power and support fast wireless communications. We chose the TelosB because it includes a 250 kbps IEEE 802.15.4 wireless transceiver; on-board temperature and light sensors; low power consumption; an on-board antenna; simple programming and collection interface; and a 10+6-pin expansion slot allowing analog to digital, UART, I2C, and digital connections [6]. This selection gives us a very flexible platform and this is what makes our system viable parking structures.

Potential sensors to detect vehicle presence include infrared range finders, pressure sensors, and inductive sensors. We chose to use the LV-MaxSonar®-EZ1™ by MaxBotix® (hereafter "EZ1") because of its price and feature set. The EZ1 supports a wide range of standard supply voltage inputs. This allows a great deal of flexibility in choosing how to power the EZ1. It can detect objects from 6 inches to 254 inches; this is enough range to cover the complete parking space. The EZ1

supports many output formats including pulse width, analog voltage, and serial digital giving us multiple ways to connect the same sensor [3].

3.2.2   Base Station

Keeping with our theme of low-cost, the requirements of the base station are minimal. We only need the features found on any low-end Unix or Linux computer. The base station requirements are the capability to connect to the Internet and a USB connection. We recommend a small, power efficient and cost effective plug computer [1].

The base station must have IEEE 802.15.4 communications with the PSMs, this can be achieved by connecting a TeloB via USB completing the bridge between the sensor network and the Internet.

3.3   Collection Software

3.3.1   Parking Space Monitors

Since we are using TinyOS on TelosB platforms, we are able to use the included mesh networking and communication protocols, as well as having native support for all the connection interfaces described earlier.

The PSMs receive configuration data from the base station and then begin monitoring sensor values. During a sampling period the PSMs reads the sensors ten times and average the values. Those values are then transmitted to the base station every 15 seconds. One of our key design goals was wireless reliability. Packets containing sensor details must always make it back to the base station for processing. This can be difficult in dynamic environment such as a parking structure, especially a multi-floored garage where line of site is impossible, and transmissions will attenuate through the walls. To overcome this, we attempted to use Collection Tree Protocol (CTP).

CTP is a dynamic point-to-sink routing protocol, which creates a tree topology for the network through which all packets travel to the base station (the "root node" or "sink"). By creating this topology, motes can build efficient routes for collecting data at a centralized point. CTP does not handle point-to-point communication, though it does provide support for broadcasting to all nodes [10]. In our final implementation, due to a miscommunication about integrating the various parts of the project, CTP was left out. However, initial tests of a non-trivial CTP network gave indications that it would function well within our project.

3.3.2   Base Station

The base station software has multiple functions. It configures the PSMs, monitors their status, collects data from them, and sends the data to the backend over the Internet.

Since the base station needs to communicate with the motes over IEEE 802.15.4, a TelosB is connected via USB as part of the base station. The TelosB runs the

default BaseStation app (included with the TinyOS install) with only minor tweaks.

The base station is able to send configuration packets to the PSMs. This configuration process allows the PSMs to be assigned parking space identifiers. The configuration packet structure was designed to be easily extended to support future needs. For example, the base station could adjust sensor read rates or data transmission rates based on the time of day to increase battery life.

The base station's primary purpose is to collect sensor readings from the PSMs and relay it to the aggregation software. Once the data is collected, it is encoded as a JSON packet (as seen in Figure 2) and transfered via HTTP PUT to the backend for processing.

3.4   Aggregation Software

Data leaving the base station is directed over the Internet to our aggregation software hosted on Google's AppEngine [2]. Using AppEngine provides us with: free hosting service for light usage with the option to pay as usage increases, a Python programming environment for rapid development, and the ability to transparently scale up as usage increases.

Scalability is important if we are to have one sensor per parking space for every parking lot. Thus our backend must be able to accept, store, and recall large amounts of current and historical data. A traditional server model may have worked for this prototype, but if the project starts scaling up we would quickly start running into barriers.

Though transparent to users of our system, this non-SQL based data store forces some differences from a SQL based data store. While queries look like SQL, they are actually "GQL," a SQL-like language. In general this does not bring any problems, but because of the organization of the data store imprecise queries are not acceptable. You must know the data you are looking for. You cannot SELECT one or two columns from a row, you must take the entire row or an identifier for the row. As such, proximity searches are difficult because there is not concept of "similar" or "like" queries. You either know how to find the data or you can't find the data at all. However, with the help of the Geomodel API [1] we can use geohashing for proximity searches. This allows the end-user to identify nearby parking areas.

3.4.1   Frontend

Our prototype frontend is an HTML interface that allows viewing information about a parking lot. The user selects a lot and our frontend provides a map of the area and marks lots within a two mile radius. It also displays a "health" indicator giving an approximate empty-to-full ratio, as well as a percentage and a list of full spaces and their usage duration.

We also provide a chart to indicate the average fullness of a lot over a time frame. For example, Figure 6 shows the average fullness of the "DUC" lot from 5:30am until

[1] http://www.wikipedia.org/wiki/Plug_computer

Figure 6: Example of a usage chart for the 5:30am–7:00am time slot.

7:00am. At the 6:06am time slot there are 11 spaces being used, however only 20 minutes earlier only 6 spaces are being used. This information allows a commuter, or one time visitor, to find that by arriving a little earlier they have a much greater chance of easily finding a parking space.

4. EXPERIMENT

Since our system was split into two major components, we were able to test both somewhat independently. The first component included physically sensing the vehicle and relaying the data to the backend [2]. The second component involved testing the backend works for extended periods of time and correctly logs/recalls historical data.

4.1 Sensing and Sending

Figure 7 shows a PSM mounted on a tripod; this set-up was used to test the system with an actual vehicle. Testing found issues with the sensing range of our PSMs. The EZ1 specification claims support for distances up to 254-inches, however we were only able to detect a vehicle up to approximately 120-inches (10 feet). Though this did not cause an issue detecting a vehicle, it could be an issue with detecting smaller objects such as motorcycles. More testing and analysis would need to be performed to determine if the ten foot range of the EZ1 would really be an issue. It was determined while testing this issue that it was caused by the supply voltage to the EZ1. The EZ1 will support a supply voltage of 2.5V - 5.5V. We chose to power the EZ1 using the TelosB analog supply which supplies about 3.1V. When we used an external supply to test just the EZ1 at 5V it had a much greater range compared to 3.1V. If after more testing it is determined to be an issue, there are solutions such as using an external power supply for the EZ1 or possibly mounting the sensor above the space to cover a greater area of the space.

4.2 Backend and Frontend

Since it was not feasible to test a large scale deployment due to cost, time, and logistical constraints a base station simulator was developed. Instead of interacting with a sensor network for gathering data, the simulator was designed to generate a number of changes to a lot every one to two minutes. The changes include

---

[2] A video is available at `http://www.youtube.com/watch?v=cE2xn1PHDlI`.

Figure 7: Parking Space Monitor on tripod for testing.

modifying the status of a random number of parking spaces with some weight attached to the likelihood of a space becoming empty for a time interval. For example, during the overnight hours there only a 16% chance of a parking space becoming full causing the overnight hours to present less lot usage.

This simulator helped test the ability of the backend to continue to be stable over a number of hours while periodically inserting new log information and maintaining the current status of a parking lot. It also helped generate large amounts of data for testing the analytics portion of our frontend (charting for time periods).

5. RELATED WORK

There are other proposals and existing systems for parking guidance and information, here we'll discuss and differentiate ourselves from these systems.

Signal-Park, uses sonar sensors above each parking spot to detect the presence of a vehicle [4]. Approximately three times per second, the sensor is queried to determine vehicle presences and updates a central computer. This system uses serial lines to connect each sensor with the central computer, leading to more costly deploy. It also appears to be limited to the parking lot that it is deployed in and lacks the ability to aggregate the information for others to benefit. It would certainly be possible to improve the connectivity of this system to integrate with our backend aggregation system to take advantage of existing installations.

Streetline, a startup in the San Francisco area, uses a wireless network to link all the sensors together [5]. The Streetline system has recently (Summer 2010) begun to roll out sensors and upgraded meters for select areas in the San Francisco area [8, 7]. This system uses surface (and sub-surface) sensor mounts that contain a wireless transmitter and a magnetometer for vehicle detection. Status information is sent to a central aggregation server where users can see the status of street parking via their web browser, text message, or smart phone; it also allows municipal workers to identify vehicles that have not paid fully. Presuming the Streetline trial is successful and the system expands, our system and their system should be able to integrate with little effort and could be seen as competitors. We view this a positive, as it demonstrates that our project is a useful idea with potential buyers available.

6. LESSONS LEARNED
If we were to pursue this project into the future we would probably modify the design to allow a single TelosB to monitor multiple sensors. The ultrasonic range finder we used has multiple output types, one of them being a two wire digital serial connection. This allows for up to twelve EZ1s to be used with a single serial connection. This allows either two sensors per space to improve accuracy and monitor six spaces per TelosB or continue with a single range finder per space and monitor twelve parking spaces per TelosB. This could greatly decrease the cost and power consumption of the overall system.

Another interesting lesson was using Google's AppEngine for the backend. The cost structure of AppEngine is set up to be free as long as resource usage stays below a daily threshold. After modifying the frontend to automatically update the lot usage, we easily surpassed the daily threshold giving us two options: back the AppEngine instance with some money to increase our limits or refactor the code to be less resource intensive. Luckily, we were able to find a few code optimizations that greatly reduced the resources for the automatic update code putting us under the daily allotment again.

7. CONCLUSIONS AND FUTURE WORK
This project created an extensible wireless sensor that is able to detect the presence of a vehicle in a parking space. That data is transferred from the sensor, to the base station, and over the Internet to our backend aggregator. Once at our backend, the data is logged and inserted into a scalable web infrastructure. At any point an end-user is able to query our web service to determine the current fullness of a lot at their destination and, if desired, find the fullness of lots that are near-by to their destination. The end-user is also able to discover when the lot is likely to be most empty or most full. Moreover, through the use of a web-based API using the standard JSON format, application developers can interact with our system hassle-free allowing rich web apps, smart-phone apps, or data mining opportunities.

As with many projects, there are both many features

that can still be added and much research into the viability of deploying such a sensor network. Feature-wise, the web API could support richer queries. For example, the ability to select multiple date ranges for charting allowing a user to view average aggregate data for the weekdays over the past month to decide if a parking lot should expand. Another option is to integrate our system with a mapping service and combine up to the minute traffic reports to allow a user to plan their trip while being mindful of the route, traffic conditions, and available parking. On the hardware side, a careful evaluation of battery life and how to increase battery life would be helpful. Stripping away unused components from the TelosB could help greatly as well as integrating solar panels for power support and creating an external casing to protect the sensitive electronics from the environment.

8. REFERENCES
[1] Geomodel.
http://code.google.com/p/geomodel/.
[2] Google AppEngine.
http://code.google.com/appengine/.
[3] LV-MaxSonar-EZ1 High Performance Sonar Range Finder. http://www.maxbotix.com/uploads/LV-MaxSonar-EZ1-Datasheet.pdf.
[4] Signal-park parking management solutions.
http://www.signalpark.com/.
[5] Streetline parking management system.
http://www.streetlinenetworks.com/.
[6] Telos (Rev B): PRELIMINARY Datasheet, December 2004. www2.ece.ohio-state.edu/~bibyk/ee582/telosMote.pdf.
[7] SF*park*, 2010. http://sfpark.org/.
[8] K. Barry. City parking smartens up with Streetline. *Wired Magazine*, November 2010. http://www.wired.com/autopia/2010/11/city-parking-smartens-up-with-streetline/.
[9] D. Crockford. JavaScript Object Notation. http://www.json.org/.
[10] R. Fonseca, O. Gnawali, K. Jamieson, S. Kim, P. Levis, and A. Woo. The collection tree protocol (ctp). TEP 123 (Documentary). TinyOS-2.x, URL http://www.tinyos.net/tinyos-2.x/doc/html/tep123.html.
[11] C. Pautasso, O. Zimmermann, and F. Leymann. Restful web services vs. "big'" web services: making the right architectural decision. In *Proceeding of the 17th international conference on World Wide Web*, WWW '08, pages 805–814, New York, NY, USA, 2008. ACM.
[12] A. Sakai, K. Mizuno, T. Sugimoto, and T. Okuda. Parking guidance and information systems. In *Vehicle Navigation and Information Systems Conference, 1995. Proceedings. In conjunction with the Pacific Rim TransTech Conference. 6th International VNIS. 'A Ride into the Future'*, pages 478 –485, July 1995.
[13] L. Vincent. Taking online maps down to street level. *Computer*, 40(12):118 –120, 2007.